

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 1 of 50

You have Azure Databricks tables and a Fabric lakehouse.

You need to create a new Fabric artifact to combine data from both architectures. The solution must use data pipelines for the Azure Databricks data and shortcuts for the existing Fabric lakehouse.

What Fabric artifact should you create?

- a data warehouse
- a Dataflow Gen2 query
- a lakehouse
- a semantic model

✓ This answer is correct.

Only Fabric lakehouses can shortcut to other lakehouses. Fabric data warehouses can use data Pipelines but cannot use shortcuts.

[OneLake shortcuts - Microsoft Fabric | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 2 of 50

You have a Fabric workspace that contains a lakehouse named Lakehouse1.

You need to create a data pipeline and ingest data into Lakehouse1 by using the Copy data activity.

Which properties on the General tab are mandatory for the activity?

- Name and Retry only
- Name and Timeout only

- Name only

✓ This answer is correct.

- Name, Timeout, and Retry

This answer is incorrect.

- Retry only
- Timeout only

Besides copy data, only the name must be defined. All the other properties are optional.

[Lakehouse tutorial - Ingest data into the lakehouse - Microsoft Fabric | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 3 of 50

You have a Fabric tenant that contains two lakehouses named Lakehouse1 and Lakehouse2. Lakehouse1 contains a table named FactSales that is partitioned by a column named CustomerID.

You need to create a shortcut to the FactSales table in Lakehouse2. The shortcut must only connect to data for CustomerID 100.

What should you do?

- Add a filter activity after the copy data activity in a data pipeline.
- As you create the shortcut select the CustomerKey=100 folder under the FactSales folder in Files.

This answer is incorrect.

- As you create the shortcut, select the CustomerKey=100 folder under the FactSales folder in Tables.

✓This answer is correct.

- In the semantic models connected to the lakehouses, add a report-level filter for CustomerKey = 100.

During the shortcut setup process, you can expand the FactSales folder to see each folder per CustomerID partition and select the folder for CustomerID=100. These folders are unavailable under Files, and all other options will connect to all the customer data in the shortcut.

[Referencing data to a Lakehouse using shortcuts - Microsoft Fabric | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 4 of 50

You have a Fabric tenant that contains a lakehouse.

You plan to use a Fabric notebook and PySpark to read sales data and save the data as a Delta table named Sales. The table must be partitioned by Sales Year and Quarter.

You load the sales data to a DataFrame named df that contains a Year column and a Quarter column.

Which command should you run next?

- `df.write.mode("overwrite").format("delta").partitionBy("Year",Quarter").save("Files/Sales")`
 - `df.write.mode("overwrite").format("delta").partitionBy("Year","Quarter").save("Tables/Sales")`
- ✓This answer is correct.
- `df.write.mode("overwrite").format("parquet").partitionBy("Year", "Quarter").save("Files/Sales")`
 - `df.write.mode("overwrite").format("parquet").partitionBy("Year", "Quarter").save("Tables/Sales")`

To save a DataFrame in the Delta format, you must use `format("delta")`. While a DataFrame can be saved to the Files section of a lakehouse, it will not be considered a table.

[Lakehouse tutorial - Prepare and transform data in the lakehouse - Microsoft Fabric | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 5 of 50

You have a Fabric workspace that contains a Microsoft Power BI report.

You need to modify the column names in the Power BI report without changing the original names in the underlying Delta table.

Which warehouse object should you create?

- columnstore index
- schema
- table-valued function
- view

✓This answer is correct.

A view provides a convenient way to encapsulate additional query logic, such as renaming columns, filtering, aggregating, etc. Views contain only a query definition and do not change the underlying tables.

[CREATE VIEW \(Transact-SQL\) - SQL Server | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 6 of 50

You have a Fabric tenant.

Your company has 1 TB of legacy accounting data stored in an Azure Data Lake Storage Gen2 account. The data is queried only once a year for a few ad-hoc reports that submit very selective queries.

You plan to create a Fabric lakehouse or warehouse to store company sales data. Developers must be able to build reports from the lakehouse or warehouse based on the sales data. The developers must also be able to do ad-hoc analysis of the legacy data at the end of each year.

You need to recommend which Fabric architecture to create and the process for integrating the accounting data into Fabric. The solution must minimize administrative effort and costs.

What should you recommend?

- Ingest the sales data into the Fabric lakehouse and set up a shortcut to the legacy accounting data in the storage account.

✓ This answer is correct.

- Ingest the sales data into the Fabric lakehouse and use a pipeline to move the legacy accounting data into the lakehouse.

- Ingest the sales data into the Fabric warehouse and use a pipeline to move the legacy accounting data into the warehouse.

This answer is incorrect.

- Set up a lakehouse with a shortcut to the legacy accounting data. Ingest the sales data into the Fabric warehouse and add the SQL analytics endpoint of the lakehouse to the warehouse for cross querying.

Since the legacy accounting data is only accessed once a year for a few ad-hoc queries that are highly selective, there is no need to move the data into a Fabric workspace. Shortcuts enable the querying of remote data without having to move

the data. Shortcuts are only supported in a Fabric lakehouse. While you can add the SQL endpoint of a lakehouse to a warehouse for cross database querying, that is not the simplest method. The simplest method is to use a shortcut.

[OneLake shortcuts - Microsoft Fabric | Microsoft Learn](#)

Next >

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 7 of 50

You have a Fabric workspace that contains a data pipeline with a fact table and two dimension tables. The fact table contains customer data. One dimension table contains customer information and a column with Customer ID information, and the other dimension table contains calendar information and a column with Date ID information.

You need to ensure that each customer's sales data is provisioned to their own Parquet file under the Parquet folder structure.

Which data pipeline configuration should you implement?

- Add a SecureString parameter for Customer ID.
- Increase the Concurrency count.
- Partition by customer ID on the customer dimension table.

This answer is incorrect.

- Partition by customer ID on the fact table.

✓This answer is correct.

Partitioning determines the Parquet file structure, depending on the column or columns selected. Partitioning the fact table by customer ID will give each customer ID its own file.

[Load data to Lakehouse using partition - Microsoft Fabric | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 8 of 50

You have a Fabric workspace that contains a Microsoft Power BI report named Sales.

You plan to use Dataflow Gen2 to add an additional column to the report. The new column must be based on the unit price of a product. Any product that has a unit price that is greater than \$1,000 must be labeled as High, while any product that has a unit price that is less than \$1,000 must be labeled as Regular.

What should you select on the Add column tab in Power Query Editor?

Duplicate column

Conditional column

✓ This answer is correct.

Index column

Merge columns

The Conditional column option enables adding new columns whose values will be based on one or more conditions applied to the existing table columns.

[Add a conditional column - Power Query | Microsoft Learn](#)

Next >

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 9 of 50

You have a Fabric workspace named Workspace1 that contains a data pipeline named Pipeline1.

You plan to use the Office 365 Outlook activity to send an email message each time Pipeline1 experiences issues with pipeline connectors.

You need to connect the Office 365 Outlook activity to each main pipeline activity. The solution must minimize the number of email messages sent by the activity.

Which connection should you use for the Office 365 Outlook activity?

On completion

On fail

✓ This answer is correct.

On skip

On success

This answer is incorrect.

On fail is correct because it will only send a notification when there is an issue caused by an activity failing. All the other options will either notify incorrectly or notify every time.

[Data pipeline runs - Microsoft Fabric | Microsoft Learn](#)

[Ingest data with Microsoft Fabric - Training | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 10 of 50

You have a Fabric tenant that contains a workspace named Workspace1. Workspace1 contains a data pipeline named Pipeline1 that runs in the US-West Azure region. Workspace1 also contains a semantic model named SemanticModel1 and a warehouse named Warehouse1.

You need to ensure that Pipeline1 runs at midnight (12:00 AM), and that the schedule is set to the UTC-0 time zone.

How should you configure the schedule for Pipeline1?

- Add a data pipeline notebook activity to convert the US West time zone to UTC-0.
- For Pipeline1, set the scheduler time zone to UTC-0.
- For SemanticModel1, set the schedule time zone to UTC-0.
- For Warehouse1, set the scheduler time zone to UTC-0.

The data pipeline artifact in the workspace has its own time zone setting that applies to only that data pipeline. This is where you need to configure the UTC time zone that will apply to only the data pipeline.

[Data pipeline runs - Microsoft Fabric | Microsoft Learn](#)

[Use Data Factory pipelines in Microsoft Fabric - Training | Microsoft Learn](#)

Next >

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 11 of 50

You have a Fabric workspace named Workspace1.

You plan to create a data pipeline to ingest data into Workspace1.

You need to ensure that the pipeline activity supports parameterization.

Which two activities support parameterization in the data pipeline UI? Each correct answer presents part of the solution.

Dataflow Gen2

KQL activity

notebooks

✓ This answer is correct.

SQL stored procedures

✓ This answer is correct.

user-defined functions

This answer is incorrect.

Only notebooks and SQL stored procedures provide a possibility to define parameters in the data pipeline UI. Dataflow Gen2 and KQL activity only require connection details, but no parameters can be supplied. User-defined functions cannot be added as an activity to a pipeline.

[Ingest Data with Dataflows Gen2 in Microsoft Fabric - Training | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 12 of 50

You are designing a dimension table named dimCustomer that will be used to analyze historical sales data by customer zip code. The table will be joined to a table named FactSales on a column named CustomerKey to report historical sales data by customer zip code.

The sales data must be reported based on the zip codes of customers at the time of the sale, not their most recent zip code.

You need to design dimCustomer to contain a fixed number of columns.

Which type of dimension should you choose for dimCustomer?

- type 0 slowly changing dimension (SCD)
- type 1 slowly changing dimension (SCD)
- type 2 slowly changing dimension (SCD)

✓This answer is correct.

- type 3 slowly changing dimension (SCD)

Type 0 SCD attributes never change and will not fit the requirement. Type 1 SCD overwrites the changes and historical analysis of data based on the zip code at the time of the sales will be impossible. Type 2 SCD will keep track of historical data by adding new records with new keys whenever an attribute changes. Type 3 SCD adds new columns to a table for attribute changes.

[Explore data load strategies - Training | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 13 of 50

You have a Fabric lakehouse named Lakehouse1 that contains a Dataflow Gen2 query.

You have an Azure SQL database that contains a type 2 slowly changing dimension database table named CustomerMaster.

CustomerMaster contains the following columns:

- Customer ID – Number
- EffectiveDate – Date
- Address – Text
- Status - Text

You plan to ingest CustomerMaster into Lakehouse1. The solution must only keep the latest record (unique) per Customer ID.

Which two applied steps should you use? Each correct answer presents part of the solution.

Keep top rows

Max Customer ID

Remove duplicates on the Customer ID column

✓ This answer is correct.

Remove duplicates on the CustomerMaster table

Sort on Customer ID, EffectiveDate

✓ This answer is correct.

Sorting CustomerID and EffectiveDate, and then removing duplicates on the Customer ID column is the only way to keep the correct latest row per customer

ID. All other options will not correctly keep the latest customer row per effective date.

[Sort columns - Power Query | Microsoft Learn](#)

[Working with duplicate values - Power Query | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 14 of 50

You have a Fabric warehouse.

You create a table named dimCustomer by using the following code.

```
CREATE TABLE dbo.Dim_Customer (
    CustomerKey VARCHAR(255) NOT NULL,
    Name VARCHAR(255) NOT NULL,
    Email VARCHAR(255) NOT NULL
);
```

You need to alter the table to add CustomerKey as a primary key.

Which command should you run?

- ALTER TABLE dbo.Dim_Customer add CONSTRAINT PK_Dim_Customer PRIMARY KEY CLUSTERED (CustomerKey) ENFORCED
- ALTER TABLE dbo.Dim_Customer add CONSTRAINT PK_Dim_Customer PRIMARY KEY CLUSTERED (CustomerKey) NOT ENFORCED
- ALTER TABLE dbo.Dim_Customer add CONSTRAINT PK_Dim_Customer PRIMARY KEY NONCLUSTERED (CustomerKey) ENFORCED
- ALTER TABLE dbo.Dim_Customer add CONSTRAINT PK_Dim_Customer PRIMARY KEY NONCLUSTERED (CustomerKey) NOT ENFORCED

✓ This answer is correct.

PRIMARY KEY is only supported when NONCLUSTERED and NOT ENFORCED are both used.

[Primary, foreign, and unique keys - Microsoft Fabric | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 15 of 50

You have an Azure SQL database that contains fact table named UnpostedSales. UnpostedSales contains unposted payments.

Each day payment records from the previous day are automatically truncated from the UnpostedSales table and replaced with today's payment records.

You need to use a Dataflow Gen2 query to import the data into either a lakehouse or a warehouse. The solution must ensure that all current and historical records are maintained.

What should you do to retain current and historical records during a refresh?

- Configure the refresh to append data for the query.

✓ This answer is correct.

- Configure the refresh to replace data for the query.

- Use Optimize to apply V-order for the query.

This answer is incorrect.

- Ensure that query folding occurs for the whole query.

Appending data for the query will add new rows each time a refresh occurs, ensuring that both historical and current records are kept and combined. All other options will not keep the historical records.

[Lakehouse Load to Delta Lake tables - Microsoft Fabric | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 16 of 50

You have a Fabric warehouse.

You have an Azure SQL database that contains two tables named ProductCategory and Product. Each table contains a column named ProductCategoryKey.

You plan to ingest the tables into the warehouse using Dataflow Gen2.

You need to merge the tables into a single table named Product. The combined table must contain all the rows from the Product table and the matching rows from the ProductCategory table.

Which join configuration should you use?

- a left anti join Product to ProductCategory
- a left anti join ProductCategory to Product
- a left outer join Product to ProductCategory

✓This answer is correct.

- a left outer join ProductCategory to Product

This answer is incorrect.

Only a left outer join from Product to Product Category will keep all the rows from Product but only matching rows from Product Category. The anti joins will only keep rows not found from the left table, in the right table, and the left outer join from ProductCategory to Product will start with the ProductCategory table and only keep matching rows from the Product table.

[Ingest Data with Dataflows Gen2 in Microsoft Fabric - Training | Microsoft Learn](#)

[Merge queries overview - Power Query | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 17 of 50

You have a Fabric warehouse that contains two tables named FactSales and dimGeography. The dimGeography table has a primary key column named GeographyKey. The FactSales table has a foreign key column named GeographyKey.

You create a Dataflow Gen2 query and add the tables as queries.

You plan to use the Diagram view to visually transform the data.

You need to join the two queries so that you retain all the rows in FactSales even if there are no matching rows for them in dimGeography.

What should you do after you select FactSales?

- Use Append queries as new transformation with FactSales as the first table and dimGeography as the second table.
- Use Merge queries as new transformation with Join kind set to Inner outer and FactSales as the left table and dimGeography as the right table.
- Use Merge queries as new transformation with join kind set to Left outer and FactSales as the left table and dimGeography as the right table.

✓ This answer is correct.

- Use Merge queries as new transformation with Join kind set to Right outer and FactSales as the left table and dimGeography as the right table.

Using the Diagram view, you can join two existing queries by using the Merge Queries as new transformation. A left join keeps all the rows from the left table even when there is no match for them in the right table. The Append queries as new transformation adds the queries to each other like a Union operation.

[Create your first Microsoft Fabric Dataflow Gen2 - Microsoft Fabric | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 18 of 50

You have a Fabric warehouse.

You have an Azure SQL database that contains a fact table named Sales and a second table named ExceptionRecords. Both tables contain a unique key column named Record ID.

You plan to ingest the Sales table into the warehouse.

You need to use Dataflow Gen2 to configure a merge type to ensure that the Sales table excludes any records found in the ExceptionRecords table, and that query folding is maintained.

Which applied steps should you use?

- Merge (inner join) applied step, and then the expand columns applied step
- Merge (left anti join) applied step, and then the expand columns applied step
- Merge (left anti join) applied step, delete the “expand columns” column
- Merge (right anti join) applied step, and then the expand columns applied step

A left anti join ensures that only rows not found in the ExceptionRecords table are loaded, and the expand columns step ensures that query folding is maintained for performance.

[Merge queries overview - Power Query | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 19 of 50

You have a Fabric workspace that contains a Microsoft Power BI report named Report1.

Your organization does not currently have an enterprise data warehouse.

You need to leverage dataflows to bring data into a Power BI semantic model. You notice that access to one of the data sources is restricted to narrow time windows.

What should you do?

- Create a linked table that will reference the data from another dataflow.

This answer is incorrect.

- Create a shared dataset that can be reused by multiple Power BI reports.
- Create a staging dataflow that will only copy the data from the source as-is.

✓This answer is correct.

- Create a transformation dataflow that will apply all the necessary data transformations.

A staging dataflow copies raw data “as-is” from the data source and can then be used as a data source for further downstream transformations. This is especially useful when access to a data source is restricted to narrow time windows and/or to a few users.

[Power BI usage scenarios: Advanced data preparation - Power BI | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 20 of 50

You have a Fabric tenant that contains a workspace named Workspace1. Workspace1 contains a lakehouse named Lakehouse1.

You open a notebook in Lakehouse1 and attach it to a Spark session.

You plan to start a new notebook in Lakehouse1 and attach it to the same Spark session. However, you notice that the New high concurrency session option is unavailable, and the only available option is Standard session.

You need to ensure that the high concurrency mode for notebooks is enabled.

Where can you check the high concurrency mode?

Fabric tenant settings

Lakehouse settings

This answer is incorrect.

Notebook properties from the Edit menu

Workspace settings

✓This answer is correct.

The high concurrency mode for Fabric notebooks is set at the workspace level. It is on by default; however, it can be turned off in scenarios where notebooks require dedicated compute resources.

[Configure high concurrency mode for notebooks - Microsoft Fabric | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 21 of 50

You have a Fabric workspace that contains a lakehouse named Lakehouse1. Lakehouse1 contains a Delta Parquet table named FactSales.

You use a Describe command to review the history of FactSales and notice that you have over 1000 versions of the table, and the retention policy is six months.

You need to reduce the size of the FactSales table and the number of files in the table.

What should you configure on the table?

- Apply V-Order under Maintenance.
- Delete the FactSales table from Lakehouse1.
- Run the OPTIMIZE command under Maintenance.

This answer is incorrect.

- Run the VACUUM command under Maintenance.

✓This answer is correct.

The VACUUM command will delete unreferenced files that are older than whatever the configured retention policy is set to. This will reduce the number of files and the storage size.

[Use table maintenance feature to manage delta tables in Fabric - Microsoft Fabric | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 22 of 50

You have a Fabric workspace that contains a lakehouse named Lakehouse1. Lakehouse1 contains a table named FactSales that is ingested by using a Dataflow Gen2 query.

There are several applied steps and transformations applied to FactSales during the ingestion process.

You notice that due to the number of Power Query transformations, there are occasional timeout issues for the dataflow.

You need to recommend a solution to prevent the timeout issues.

You have already confirmed that the query cannot be further optimized and that changing the refresh time does not improve the timeout issues.

Which additional action should you recommend?

- Change the FactSales destination to a newly created Fabric warehouse.
- Create a new group for the FactSales query.
- Create a second dataflow that ingests the FactSales table with no additional transformations, and then connect the original dataflow to transform the FactSales data by using this second dataflow.

✓ This answer is correct.

- Create a shortcut to FactSales, connect the dataflow to the shortcut, and then apply the transformations.

It is considered best practice to create a staging first dataflow that ingests the raw data first, and then a second dataflow to transform the data, commonly applied when there are performance or timeout issues for a query.

Best practices for designing and developing complex dataflows - Power Query | Microsoft Learn

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 23 of 50

You are planning a Fabric analytics solution.

You need to recommend a licensing strategy to support 10 Microsoft Power BI report authors and 600 report consumers.

The solution must use Dataflow Gen2 for data ingestion and minimize costs.

Which Fabric license type should you recommend?

- F16
- F32
- F64

✓ This answer is correct.

- Premium Per User (PPU)

While F32 and F16 license types will provide all the necessary set of features, these licenses are not cost-efficient because report consumers require a Pro or PPU license. Starting with the F64 license, report consumers can use a free per-user license. PPU is incorrect, because you cannot create non-Power BI items (in this case Dataflow Gen2) with PPU.

[Microsoft Fabric concepts - Microsoft Fabric | Microsoft Learn](#)

[Implement a Lakehouse with Microsoft Fabric - Training | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 24 of 50

You are planning a Fabric analytics solution for the following users:

- 2,000 Microsoft Power BI consumers without an individual Power BI Pro license.
- 32 Power BI modelers with an individual Power BI Pro license.
- 16 data scientists

You need to recommend a Fabric capacity SKU. The solution must minimize costs.

What should you recommend?

- F2
- F2048
- F32
- F64

✓ This answer is correct.

F64 is the smallest Fabric capacity (equivalent to a P1 Power BI Premium capacity) that supports premium Fabric workspaces and does not require Power BI report consumers to have individual Power BI Pro licenses. F2 and F32 are incorrect since they require that the 2,000 employees have individual Power BI Pro licenses to consume Power BI content. F2048 is incorrect since it is not the smallest capacity that meets the stated requirements.

[Microsoft Fabric concepts - Microsoft Fabric | Microsoft Learn](#)

[Implement a Lakehouse with Microsoft Fabric - Training | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 25 of 50

You have a Fabric tenant that contains a workspace named Workspace1. Workspace1 contains a lakehouse, a data pipeline, a notebook, and several Microsoft Power BI reports.

A user named User1 plans to use SQL to access the lakehouse to analyze data. User1 must have the following access:

- User1 must have read-only access to the lakehouse.
- User1 must **NOT** be able to access the rest of the items in Workspace1.
- User1 must **NOT** be able to use Spark to query the underlying files in the lakehouse.

You need to configure access for User1.

What should you do?

- Add User1 to the workspace as a member, share the lakehouse with User1, and select **Read all SQL Endpoint data**.
- Add User1 to the workspace as a viewer, share the lakehouse with User1, and select **Read all SQL Endpoint data**.
- Share the lakehouse with User1 directly and select **Build reports on the default dataset**.
- Share the lakehouse with User1 directly and select **Read all SQL Endpoint data**.

✓**This answer is correct.**

Since the user only needs access to the lakehouse and not the other items in the workspace, you should share the lakehouse directly and select Read all SQL Endpoint data. The user should not be added as a member of the workspace. All members of the workspace, even viewers, will be able to open all Power BI reports in the workspace. The SQL analytics endpoint itself cannot be shared directly; the Share options only show for the lakehouse.

Lakehouse sharing and permission management - Microsoft Fabric | Microsoft Learn

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 26 of 50

You have a Fabric tenant.

You notice a Fabric compute usage issue, which is causing performance issues.

You need to increase the Fabric capacity unit size.

What should you use?

Azure portal

✓ This answer is correct.

Fabric admin portal

Microsoft Power BI settings

Microsoft Purview hub

The Azure Fabric capacity portal is the only location where the capacity size can be configured or changed.

[Scale your Fabric capacity - Microsoft Fabric | Microsoft Learn](#)

[Configure and manage capacities in Power BI Premium - Power BI | Microsoft Learn](#)

Next >

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 27 of 50

You have a new Fabric tenant.

You need to recommend a workspace architecture to meet best practices for content distribution and data governance.

Which two actions should you recommend? Each correct answer presents part of the solution.

Create a copy of each semantic model in each workspace.

Create direct query semantic models in each workspace.

This answer is incorrect.

Place semantic models and reports in separate workspaces.

✓This answer is correct.

Place semantic models and reports in the same workspace.

This answer is incorrect.

Reuse shared semantic models for multiple reports.

✓This answer is correct.

Using shared semantic models for multiple reports enables the reusability of items, while placing semantic models and reports in separate workspaces ensures that the data governance recommended practices are in place.

[Create reusable Power BI assets - Training | Microsoft Learn](#)

[Manage the analytics development lifecycle - Training | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 28 of 50

You have a Fabric tenant that has XMLA Endpoint set to Read Write.

You need to use the XMLA endpoint to deploy changes to only one table from the data model.

What is the main limitation of using XMLA endpoints for the Microsoft Power BI deployment process?

- A PBIX file cannot be downloaded from the Power BI service.

✓ This answer is correct.

- Only the user that deployed the report can make changes.
- Table partitioning is impossible.
- You cannot use parameters for incremental refresh.

Whenever the semantic model is deployed/changed by using XMLA endpoints, there is no possibility to download the PBIX file from the Power BI service. This means that no one can download the PBIX file (even the user who deployed the report). Table partitioning, as well as using parameters, is still supported, thus doesn't represent a limitation.

[Manage a Power BI dataset using XMLA endpoint - Training | Microsoft Learn](#)
[Semantic model connectivity and management with the XMLA endpoint in Power BI - Power BI | Microsoft Learn](#)
[Manage the analytics development lifecycle - Training | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 29 of 50

You have a Parquet file named `Customers.parquet` uploaded to the Files section of a Fabric lakehouse.

You plan to use Data Wrangler to view basic summary statistics for the data before you load it to a Delta table.

You open a notebook in the lakehouse.

You need to load the data to a pandas DataFrame.

Which PySpark code should you run to complete the task?

`df = pandas.read_parquet("/lakehouse/default/Files/Customers.parquet")`

`df = pandas.read_parquet("/lakehouse/Files/Customers.parquet")`

This answer is incorrect.

`import pandas as pd df = pd.read_parquet("/lakehouse/default/Files/Customers.parquet")`

✓ This answer is correct.

`import pandas as pd df = pd.read_parquet("/lakehouse/Files/Customers.parquet")`

To load data to a pandas DataFrame, you must first import the pandas library by running `import pandas as pd`. Pandas DataFrames use the File API Path vs. the File relative path that Spark uses. The File API Path has the format of `lakehouse/default/Files/Customers.parquet`.

[Accelerate data prep with Data Wrangler - Microsoft Fabric | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 30 of 50

You have a Fabric lakehouse that contains a managed Delta table named Product.

You plan to analyze the data by using a Fabric notebook and PySpark.

You load the data to a DataFrame by running the following code.

```
df = spark.sql("SELECT * FROM Product")
```

You need to display the top 100 rows from the DataFrame.

Which PySpark command should you run?

- `describe(df.limit(100))`
- `df.describe(100)`
- `df.printSchema(100)`
- `display(df.limit(100))`

✓ This answer is correct.

The `display` PySpark method is used to display data in a DataFrame. To limit the data displayed, `limit(100)` can be specified.

[Work with data in a Spark dataframe - Training | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 31 of 50

You are profiling the data stored in a Fabric lakehouse.

You run the following statement.

```
df.describe().show()
```

Which three functions will be included in the results for the numeric data? Each correct answer presents a complete solution.

AVG

This answer is incorrect.

COUNT

✓ This answer is correct.

DISTINCTCOUNT

MEAN

✓ This answer is correct.

STD (standard deviation)

✓ This answer is correct.

TOP

`describe` is used to generate descriptive statistics of the DataFrame. For numeric data, results include `COUNT`, `MEAN`, `STD`, `MIN`, and `MAX`, while for object data it will also include `TOP`, `UNIQUE`, and `FREQ`.

[Explore and transform data in a lakehouse - Training | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 32 of 50

You have a Fabric tenant that contains a workspace named Workspace1. Workspace1 contains a lakehouse. The lakehouse contains a table named Customers and a Fabric notebook.

You plan to use the notebook to profile the data.

In the notebook, you set up the following DataFrame that references the Customers table.

```
df = Spark.sql("Select * from Customers")
```

You need to profile the data in the DataFrame. The solution must minimize administrative effort.

What should you do?

- Create a pandas DataFrame first, and then visualize the data in an embedded Microsoft Power BI report by running QuickVisualize method on the pandas DataFrame.
- Display the DataFrame by running `display(df)`, and then clicking the **Inspect** button.

✓ This answer is correct.

- Display the DataDrame by running `display(df)`, and then switching to chart view.
- Visualize the data in an embedded Microsoft Power BI report by running the QuickVisualize method on the DataFrame.

While you can embed a Power BI report in a notebook and analyze the data by using the QuickVisualize method, this does not generate the statistics required for profiling. The Chart option is a powerful tool to visualize data, but it does not provide statistics. The Inspect button provides summarized data distribution and statistics for each column.

[Notebook visualization - Microsoft Fabric | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 33 of 50

You have a Microsoft Power BI report that contains two metrics named starting balance and ending balance. Additional fields are used for the customer segment.

You need to recommend a visual to display the value breakdown between the starting and ending balance of each customer segment.

Which visual should you recommend?

- ribbon chart
- stacked bar chart
- treemap
- waterfall chart

✓ This answer is correct.

The waterfall chart is designed to map the value breakdown between a starting and end value, illustrating the plus or minus value changes for each breakdown category. Therefore, this is the most efficient visual for this scenario.

[Waterfall charts in Power BI - Power BI | Microsoft Learn](#)

Next >

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 34 of 50

You have Fabric tenant that contains a workspace named Workspace1. Workspace1 contains a lakehouse names Lakehouse1. Lakehouse1 contains a dimension table called dimension_city. The table contains a column named City and a column named SalesTerritory.

You need to visualize the number of cities in each sales territory in a bar chart. The sales territory must be on the X axis and the number of cities on the Y axis.

You begin to create the following PySpark code in a Fabric notebook attached to Lakehouse1.

```
from matplotlib import pyplot as plt

# Get the data as a Pandas dataframe
data = spark.sql("SELECT SalesTerritory, COUNT(City) AS CityCount \
                  FROM dimension_city \
                  GROUP BY SalesTerritory \
                  ORDER BY SalesTerritory").toPandas()

# Clear the plot area
plt.clf()

# Create a Figure
fig = plt.figure(figsize=(12,8))
```

You need to complete the code to meet the analysis requirements.

How should you complete the code?

- fig.bar(x=data['SalesTerritory'],
height=data['CityCount'])fig.xlabel('SalesTerritory') fig.ylabel('Cities')
fig.show()
- plt.bar(x=data['CityCount'], height=data['SalesTerritory'])
plt.xlabel('SalesTerritory') plt.ylabel('Cities') plt.show()

plt.bar(x=data['SalesTerritory'], height=data['CityCount'])
plt.xlabel('SalesTerritory') plt.ylabel('Cities') plt.display()

This answer is incorrect.

plt.bar(x=data['SalesTerritory'], height=data['CityCount'])
plt.xlabel('SalesTerritory') plt.ylabel('Cities') plt.show()

✓ This answer is correct.

After you create a figure, you must create a bar plot of city counts by SalesTerritory by running the following code.

```
plt.bar(x=data['SalesTerritory'], height=data['CityCount'])

plt.xlabel('SalesTerritory')
plt.ylabel('Cities')

plt.show()
```

[Visualize data in a Spark notebook - Training | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 35 of 50

You have a table named Sales that contains the following columns:

- Order_ID
- Customer_ID
- Product_ID
- Quantity
- Sales_Date

You need to write a SQL statement to find the total number of products sold for each Product_ID in January 2024.

What should you run?

- `SELECT Product_ID, COUNT(*) FROM Sales WHERE Sales_Date BETWEEN '2024-01-01' AND '2024-01-31' GROUP BY Product_ID`
- `SELECT Product_ID, COUNT(Quantity) FROM Sales WHERE Sales_Date LIKE '2023-01-%' GROUP BY Product_ID`
- `SELECT Product_ID, SUM(Quantity) FROM Sales WHERE MONTH(Sales_Date) = 1 AND YEAR(Sales_Date) = 2024 GROUP BY Product_ID`

✓ This answer is correct.

- `SELECT Product_ID, SUM(Quantity) FROM Sales WHERE Sales_Date > '2024-01-01' AND Sales_Date < '2023-01-31' GROUP BY Product_ID`

The following SQL statement is the only one that correctly filters to sales for January 2024 for each product, with a total sum of the sales amount.

```
SELECT Product_ID, SUM(Quantity)  
FROM Sales
```

```
WHERE MONTH(Sales_Date) = 1 AND YEAR(Sales_Date) = 2024  
GROUP BY Product_ID
```

[WHERE \(Transact-SQL\) - SQL Server | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 36 of 50

You have a Fabric tenant that contains a workspace named Workspace1. Workspace 1 contains a lakehouse named Lakehouse1.

You plan to use Microsoft SQL Server Management Studio (SSMS) to write SQL queries against Lakehouse1.

Where can you find the SQL connection string for Lakehouse1?

- in the Lakehouse settings under Copy SQL connection string

✓ This answer is correct.

- in the Tenant settings under Microsoft Fabric
- in the Workspace settings under Azure connections
- in the Workspace settings under Data connections

The connection string for the SQL endpoint can be found under the Lakehouse settings under the SQL analytics endpoint.

[Connectivity to data warehousing - Microsoft Fabric | Microsoft Learn](#)

Next >

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 37 of 50

You have a Fabric warehouse that contains a table named customer_info. The table contains the following columns:

- customer_id
- name
- email
- join_date

The warehouse also contains a table named order_info. The table contains the following columns:

- order_id
- customer_id
- order_total
- order_date

You need to write a SQL query that returns all the customers who have not yet placed an order (purchased).

Which SQL query should you run?

`SELECT customer_info.name FROM customer_info JOIN order_info ON customer_info.customer_id = order_info.customer_id WHERE order_info.order_id IS NOT NULL`

`SELECT name FROM customer_info INNER JOIN order_info ON customer_info.customer_id = order_info.customer_id WHERE order_info.order_id IS NULL`

This answer is incorrect.

`SELECT name FROM customer_info LEFT JOIN order_info ON customer_info.customer_id = order_info.customer_id WHERE order_info.order_id IS NULL`

✓This answer is correct.

SELECT name FROM customer_info WHERE customer_id IN (SELECT customer_id FROM order_info)

The following SQL statement only returns a list of customers for whom the customer_ID was not found in the order_info table. The WHERE clause of order_info is NULL, which means that the customers have no orders found in the order_info table.

```
SELECT name  
FROM customer_info  
LEFT JOIN order_info ON customer_info.customer_id =  
order_info.customer_id  
WHERE order_info.order_id IS NULL
```

The two following SQL statements only return customers who made purchases.

```
SELECT name  
FROM customer_info  
INNER JOIN order_info ON customer_info.customer_id =  
order_info.customer_id  
WHERE order_info.order_id IS NULL
```



```
SELECT name  
FROM customer_info  
WHERE customer_id IN  
( SELECT customer_id  
    FROM order_info )
```

This following SQL statement returns the names of customers who made purchases.

```
SELECT customer_info.name  
FROM customer_info  
JOIN order_info ON customer_info.customer_id = order_info.customer_id  
WHERE order_info.order_id IS NOT NULL
```

[Joins \(SQL Server\) - SQL Server | Microsoft Learn](#)

Next >

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 38 of 50

You have a Fabric warehouse.

You write the following T-SQL statement to retrieve data from a table named Sales to display the highest sales amount for specific customers.

```
SELECT TOP 10 CustomerKey  
 , SalesAmount  
 , ROW_NUMBER() OVER(ORDER BY SalesAmount DESC) AS Ranking  
FROM dbo.Sales  
WHERE CustomerKey IN (1, 2, 3)
```

The following is an example of the expected result.

CustomerKey	SalesAmount	Ranking
1	100	<value1>
2	100	<value2>

What is the Ranking value (value1 and value2) for the first two records?

0 and 1

This answer is incorrect.

1 and 1

1 and 2

✓This answer is correct.

10 and 9

The `ROW_NUMBER()` function simply returns the sequential number of a row within a partition, starting at 1 for the first row in each partition. Ties are not considered in the evaluation.

[ROW_NUMBER \(Transact-SQL\) - SQL Server | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 39 of 50

You have a Fabric tenant that contains a workspace named Workspace1. Workspace1 is assigned to an F64 capacity and contains a lakehouse. The lakehouse contains one billion historical sales records and receives up to 10,000 new or updated sales records throughout the day at 15-minute intervals.

You plan to build a custom Microsoft Power BI semantic model and Power BI reports from the data. The solution must provide the best report performance while supporting near-real-time (NRT) data reporting.

Which Power BI semantic model storage mode should you use?

Direct Lake

✓ This answer is correct.

Import and Direct Lake combined

DirectQuery

This answer is incorrect.

Import

Direct Lake storage mode provides NRT access to data, while providing performance close to Import storage mode and much better performance than DirectQuery. DirectQuery provides NRT access to data, but queries can run slowly when working with large datasets. Import produces fast performance; however, it requires data to be loaded to the memory of Power BI and will not provide NRT. Direct Lake tables cannot currently be mixed with other table types, such as Import, DirectQuery, or Dual, in the same model. Composite models are not yet supported.

[Learn about Direct Lake in Power BI and Microsoft Fabric - Power BI | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 40 of 50

You have a Fabric workspace that contains a lakehouse named Lakehouse1.

Lakehouse1 requires additional time intelligence calculations added to its semantic model. The model has XMLA read/write permissions enabled.

You need to add a calculation group to the Lakehouse1 semantic model.

What should you use?

a Fabric notebook

DAX Studio

This answer is incorrect.

Lakehouse explorer

Tabular Editor

✓This answer is correct.

Only Tabular Editor 2/3 can currently add calculation groups to a lakehouse semantic model.

[Create calculation groups - Training | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 41 of 50

You are developing a complex semantic model that contains more than 20 date columns.

You need to conform the date format for all the columns as quickly as possible.

What should you use?

- ALM Toolkit
- DAX Studio
- Tabular Editor
- VertiPaq Analyzer

✓ This answer is correct.

Tabular Editor supports semantic model modifications and saving back these changes to the semantic model, whereas DAX Studio and VertiPaq Analyzer support only read-only operations. ALM Toolkit is used for schema comparison for semantic models.

[External tools in Power BI Desktop - Power BI | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 42 of 50

You have a Fabric workspace and a Microsoft Power BI semantic model that contains the following tables:

- Sales (ProductKey,CustomerKey,DateKey,SalesAmount)
- Product (ProductKey,ProductName,ProductCategory)
- Date (DateKey,Date,Month,Year)
- Customer (CustomerKey,CustomerName,CustomerCity)

The Product table has a 1-to-many relationship with the Sales table based on ProductKey.

The Customer table has a 1-to-many relationship with the Sales table based on CustomerKey.

The Date table has a 1-to-many relationship with the Sales table based on DateKey.

You need to create a Power BI report so that end users can use a one column chart to analyze SalesAmount by ProductCategory or Year or CustomerCity. The solution must minimize development effort.

What should you do?

- Add three bar chart visuals to the report, one by each ProductCategory, Year, and CustomerCity. Overlay the charts and use buttons and bookmarks to display one at a time.
- Create a custom visual that has custom buttons with ProductCategory, Year, and CustomerCity.
- Set up a Fields parameter with ProductCategory, Year, and CustomerCity. Use the Fields parameter in the visual.

✓ This answer is correct.

- Set up three identical report pages, each with a bar chart by either ProductCategory, Year, and CustomerCity. Set up buttons and bookmarks to navigate between the pages.

While you can use bookmarks to navigate between report pages or change the visibility of visuals, using the Fields parameter is a much easier and more efficient way of allowing an end-user to change the fields on a visual. Developing a custom visual with built-in buttons to switch the items on the axis, involves extra development effort.

[Use parameters to visualize variables - Power BI | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 43 of 50

You have a Microsoft Power BI report that contains a table visual. The visual contains three DAX measures named Sales, Units, and Customers.

You need to apply logic-based DAX formatting to the Sales measure. The solution must minimize administrative effort and prevent the modification of the other two measures.

How should you apply the logic?

Use calculation group measure formatting.

Use conditional formatting.

This answer is incorrect.

Use dynamic measure formatting.

✓This answer is correct.

Use the fields parameter.

Dynamic measure formatting is the simplest and most effective way to add logic-based formatting to a single measure. Calculation groups can add logic-based formatting, but these are applied at the visual, page, or report level, and cannot be easily added to single measures.

[Create dynamic format strings for measures in Power BI Desktop - Power BI | Microsoft Learn](#)

Next >

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 44 of 50

You have a Microsoft Power BI report that contains a bar chart visual.

You need to ensure that users can change the y-axis category of the bar chart by using a slicer selection.

Which Power BI feature should you add?

calculation groups

This answer is incorrect.

drillthrough

field parameters

✓This answer is correct.

WhatIf parameters

Field parameters allow users to change between columns that can be used on the categorical axis of visuals. All other options do not grant this ability

[Let report readers use field parameters to change visuals \(preview\) - Power BI | Microsoft Learn](#)

Next >

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 45 of 50

You are designing a Microsoft Power BI semantic model that will contain 50 different measures, such as Sales Amount, Order Quantity, and Refund Amount.

For each measure, you need to create the same set of time intelligence calculations such as month-to-date, year-to-date, and year-over-year change.

The solution must minimize administrative effort.

What should you do?

- Create a calculation group.

✓ This answer is correct.

- Create a measure folder in the report model.
- Create all the measures in the Power BI Desktop Model view.
- Use Tabular Editor to create the measures in bulk.

Calculation groups are an efficient way to reduce the number of measures in the semantic model by grouping common measure expressions. The main benefit of using calculation groups is to reduce the overall number of measures that must be created and maintained.

[Create calculation groups - Training | Microsoft Learn](#)

Next >

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 47 of 50

You have an Azure SQL database.

You have a Microsoft Power BI report connected to a semantic model that uses a DirectQuery connection to the database.

You need to reduce the number of queries sent to the database when a user is interacting with the report by using filters and/or slicers.

What should you do?

- Add apply buttons to all the basic filters.

✓ This answer is correct.

- Add Top N filters to all the visuals.
- Change default visual interaction from cross highlighting to cross filtering.
- Enable automatic page refresh for each report page.

Adding apply buttons will pause all requests to the Azure SQL database until you finalize your filter and/or slicer selections. Then a single request can be sent once the apply button is selected. The other options will not change the number of unique query requests sent to the database.

[Create Apply all and Clear all slicers buttons in reports - Power BI | Microsoft Learn](#)
[DirectQuery optimization scenarios with the Optimize ribbon in Power BI Desktop - Power BI | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 46 of 50

You have a Fabric tenant that contains a workspace named Workspace1. Workspace 1 contains a warehouse named Warehouse1. Warehouse1 contains a table named Orders that contains 20 years of historical order data.

You create a Microsoft Power BI semantic model from the Orders table.

In Power BI Desktop, you enable incremental refresh for the table and load only one week's worth of data. You publish the semantic model to Workspace1.

Due to the size of the semantic model, you need to bootstrap the initial full load.

What can you use to create the partitions in the Power BI service without processing them?

- DAX Studio
- Microsoft Visual Studio Enterprise
- Microsoft Visual Studio Code

This answer is incorrect.

- Tabular Editor

✓This answer is correct.

You can use Tabular Editor to run a Apply Refresh Policy command on a table that has an incremental refresh policy defined in Power BI Desktop. This will create the partitions based on the policy but does not process them. This method is useful when working with very large datasets where the initial full load can take many hours.

[Advanced incremental refresh and real-time data with the XMLA endpoint in Power BI - Power BI | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 48 of 50

You have a Microsoft Power BI report page that takes longer than expected to display all its visuals.

You need to identify which report element consumes most of the rendering time. The solution must minimize administrative effort and how long it takes to capture the rendering information of each element on the report page.

What should you use?

- DAX Studio
- Performance analyzer
- SQL Server Profiler
- Tabular Editor

✓ **This answer is correct.**

Performance analyzer is a built-in feature in Power BI Desktop that captures the performance information of each element on the report page. DAX Studio can be used to analyze DAX queries, whereas Tabular Editor does not capture the rendering time at all.

[Use Performance Analyzer to examine report element performance in Power BI Desktop - Power BI | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 49 of 50

You are developing a large Microsoft Power BI semantic model that will contain a fact table. The table will contain 400 million rows.

You plan to leverage user-defined aggregations to speed up the performance of the most frequently run queries.

You need to confirm that the queries are mapped to aggregated data in the tables.

Which two tools should you use? Each correct answer presents part of the solution.

DAX Studio

✓ This answer is correct.

Performance analyzer

SQL Server Profiler

✓ This answer is correct.

Tabular Editor

This answer is incorrect.

VertiPaq Analyzer

SQL Server Profiler and DAX Studio can detect whether queries were returned from the in-memory cache storage engine or pushed by DirectQuery to the data source.

[User-defined aggregations - Power BI | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 50 of 50

You have a Fabric tenant that contains a workspace named Workspace1. Workspace1 contains a warehouse that has a table named Orders.

You have a Microsoft Power BI semantic model in Power BI Desktop that sources data from the Orders table.

You need to enable incremental refresh for the table.

Which two parameters should you create?

- FullLoad and IncrementalLoad
- LowerBound and UpperBound
- RangeStart and RangeEnd

✓ This answer is correct.

- StartDate and EndDate

Incremental refresh looks for the following two parameters that are reserved keywords and case sensitive: RangeStart and RangeEnd.

[Configure incremental refresh and real-time data for Power BI semantic models - Power BI | Microsoft Learn](#)

[Submit >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 1 of 50

You are developing a Microsoft Power BI semantic model.

Two tables in the data model are not connected in a physical relationship.

You need to establish a virtual relationship between the tables.

Which DAX function should you use?

CROSSFILTER()

PATH()

TREATAS()

✓ This answer is correct.

USERELATIONSHIP()

This answer is incorrect.

TREATAS() applies the result of a table expression as filters to columns from an unrelated table. USERELATIONSHIP() activates different physical relationships between tables during a query execution. CROSSFILTER() defines the cross filtering direction of a physical relationship. PATH() returns a string of all the members in the column hierarchy.

[TREATAS function - DAX | Microsoft Learn](#)

Next >

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 7 of 50

You have a Fabric tenant that contains a workspace named Workspace1. Workspace1 is assigned to an F64 Fabric capacity and contains a warehouse.

You are working on a custom Microsoft Power BI semantic model that sources data from the warehouse tables. You apply optimization best practices to reduce the model size. You estimate that once the model is published to the Power BI service and fully loaded, it will approach 50 GB.

Which option should you configure for the semantic model from to enable the semantic model to refresh in the Power BI service?

- Large semantic model storage format

✓ This answer is correct.

- Parameters
- Query Scale-out
- Scheduled refresh

The large semantic model storage format can be enabled from the Power BI service from the semantic model settings. It will allow data to grow beyond the 10-GB limit for Power BI premium capacities or Fabric capacities of F64 or higher.

The other options do not change the default limit of 10 GB after compression.

[Configure large datasets - Training | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 34 of 50

You are designing a semantic model for a Microsoft Power BI report. You have a table named Employee that contains the following columns: EmployeeID, EmployeeName, and EmployeePosition. You have a table named Contract that contains the following columns: EmployeeID and ContractType.

You plan to denormalize both tables and include the `ContractType` attribute.

You need to ensure that all the rows in the Employee table are preserved and include any matching rows from the Contract table.

Which type of join should you specify in the Power Query Merge queries transformation?

- cross
- inner
- Left outer

✓This answer is correct.

- Left Anti Join

A left outer join keeps all the rows from the left table (Employee) and brings any matching rows from the right table (Contract). A Left Anti Join will keep only rows from the left table and exclude any matching rows from the right table. An inner join brings only matching rows from both the left and right tables, while a cross join returns the Cartesian product of the rows in both tables.

[Left outer join - Power Query | Microsoft Learn](#)

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 38 of 50

You have a Fabric tenant that contains a lakehouse named Lakehouse1.

A SELECT query from a managed Delta table in Lakehouse1 takes longer than expected to complete. The table receives new records daily and must keep change history for seven days.

You notice that the table contains 1,000 Parquet files that are each 1 MB.

You need to improve query performance and reduce storage costs.

What should you do from Lakehouse explorer?

- Manually delete any files that have a creation date that is older than seven days.
- Select **Maintenance** and run the `OPTIMIZE` command.
- Select **Maintenance** and run the `OPTIMIZE` command as well as the `VACUUM` command with a retention policy of seven days.

✓**This answer is correct.**

- Select **Maintenance** and run the `VACUUM` command with a retention policy of seven days.

The ideal file size for Fabric engines is between 128 MB and 1 GB. This improves query performance since it reduces the need to scan numerous small files. `OPTIMIZE` compacts and rewrites the files into fewer larger files. `VACUUM` removes older Parquet files that are no longer in use. While this reduces the storage size, it by itself does not reduce the number of active files that must be scanned.

[Use table maintenance feature to manage delta tables in Fabric - Microsoft Fabric](#)
| Microsoft Learn

[Next >](#)[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 40 of 50

You have a Fabric warehouse named Warehouse1.

You discover a SQL query that performs poorly, and you notice that table statistics are out of date.

You need to manually update the statistics to improve query performance for the table.

Which column statistics should you update?

- columns of the VARCHAR and NVARCHAR data type
- columns used in GROUP BY clauses

✓This answer is correct.

- columns with high cardinality
- primary key columns

When manually creating/updating statistics for optimizing query performance, you should focus on columns used in JOIN, ORDER BY, and GROUP BY clauses.

[Statistics - Microsoft Fabric | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 41 of 50

You have a Fabric tenant that contains a lakehouse.

You are creating a notebook to explore the data in the lakehouse.

You need create a query to find the total number of records in the fact table for every individual product. The displayed results must be sorted in descending order.

How should you structure the query?

`df.groupBy("ProductKey").count().sort("count", descending=True).show()`

✓ This answer is correct.

`df.groupBy("ProductKey").count().sort("ProductKey", descending=True).show()`

This answer is incorrect.

`df.groupBy("ProductKey").count().sort("count").show()`

`df.groupBy("ProductKey").count().sort("ProductKey").show()`

GroupBy will group the data per ProductKey, and then count will return the total number of records for each ProductKey. Next, we sort the data per total number of records in descending order, and finally display the DataFrame results.

[Use Apache Spark in Microsoft Fabric - Training | Microsoft Learn](#)

Next >

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 43 of 50

You have a Microsoft Power BI visual.

You use DAX query view to review the following code extracted from the visual.

```
DEFINE
    VAR __DS0Core =
        SUMMARIZECOLUMNS('Company'[Manufacturer], "Sales", 'Metrics'[Sales],
    "ConstantSales", 'Metrics'[ConstantSales])

    VAR __DS0PrimaryWindowed =
        TOPN(1001, __DS0Core, [Sales], 0, 'Company'[Manufacturer], 1)

EVALUATE
    ROW(
        "AnalyticsLine", MEDIANX(KEEPFILTERS(__DS0Core), [Sales])
    )

EVALUATE
    __DS0PrimaryWindowed

ORDER BY
    [Sales] DESC, 'Company'[Manufacturer]
```

You need to identify which type of analytics line was added to the visual.

What should you identify?

constant line

max line

median line

✓ This answer is correct.

percentile line

In the code the section containing `MEDIANX(KEEPFILTERS(__DS0Core), [Sales])` indicates that there is a median line used in the visual.

[Use Performance Analyzer to examine report element performance in Power BI Desktop - Power BI | Microsoft Learn](#)

[DAX query view in Power BI Desktop - Power BI | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 44 of 50

You use Microsoft Power BI Desktop to connect to data stored in a CSV file.

You need to use Power Query Editor to identify the percentage of valid records in a column before loading the data to a report.

Which Power Query option should you use?

Column distribution

Column profile

This answer is incorrect.

Column quality

✓This answer is correct.

Filter by value

A percentage of valid records in the column is displayed when you enable Column quality. Column distribution provides an overview of the value frequency and distribution in a column. Column profile provides statistical data about values in a column.

[Profile data in Power BI - Training | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 45 of 50

You have a Fabric lakehouse named Lakehouse1.

You use a notebook in Lakehouse1 to explore customer data.

You need to identify the rows of a DataFrame named `df_customers` in which any of the columns (axis 1 of the DataFrame) are NULL.

Which statement should you run?

`df_customers[df_customers.isnull().any(axis=1)]`

✓ This answer is correct.

`df_customers[df_customers.isnull().query(axis=1)]`

This answer is incorrect.

`df_customers[df_customers.isnull().sum(axis=1)]`

`df_customers[df_customers.nullif().any(axis=1)]`

The `isnull()` method identifies which individual values are NULL. To see these individual values in context, you should filter the DataFrame to include only rows in which any of the columns (axis 1 of the DataFrame) are NULL.

[Exercise - Explore data with NumPy and Pandas - Training | Microsoft Learn](#)

Next >

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 47 of 50

You have a Fabric warehouse.

You are writing a T-SQL statement to retrieve data from a table named Sales to display the highest sales amount for specific customers.

```
SELECT CustomerKey  
    , SalesAmount  
    , **<target1>** OVER(ORDER BY SalesAmount DESC) AS Ranking  
FROM dbo.Sales  
WHERE CustomerKey IN (1, 2, 3)
```

You need to ensure that after ties for SalesAmount, the next Sales amount increments the Ranking value by one.

The following is an example of the expected result.

CustomerKey	SalesAmount	Ranking
1	100	1
2	100	1
1	80	2

Which function should you use for <target1> in the T-SQL statement?

DENSE_RANK()

✓ This answer is correct.

NTILE()

RANK()

- `ROW_NUMBER()`

`DENSE_RANK()` function returns the rank of each row within the result set partition, with no gaps in the ranking values. The `RANK()` function includes gaps in the ranking.

[Ranking Functions \(Transact-SQL\) - SQL Server | Microsoft Learn](#)

[Next >](#)

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 49 of 50

You have an Azure SQL database that contains a table named `inventory`. The `inventory` table contains the following columns:

- `item_id`
- `category`
- `stock_quantity`
- `last_stocked_date`

You need to write a SQL statement that retrieves the latest `last_stocked_date` for each category, for which `stock_quantity` is less than 50.

Which SQL statement should you run?

`SELECT category, last_stocked_date FROM inventory GROUP BY category HAVING MIN(stock_quantity) < 50 ORDER BY last_stocked_date DESC`

`SELECT category, MAX(last_stocked_date) FROM inventory WHERE stock_quantity < 50 GROUP BY category`

✓ This answer is correct.

`SELECT category, last_stocked_date FROM inventory WHERE stock_quantity < 50 ORDER BY last_stocked_date DESC LIMIT 1`

This answer is incorrect.

`SELECT DISTINCT category, last_stocked_date FROM inventory WHERE stock_quantity < 50 ORDER BY last_stocked_date DESC`

The following SQL statement correctly returns the latest stock date per category for which the quantity is less than 50.

```
SELECT category, MAX(last_stocked_date)  
FROM inventory
```

```
WHERE stock_quantity < 50  
GROUP BY category
```

This following SQL statement does not aggregate the last stock date and will return all dates per category.

```
SELECT DISTINCT category, last_stocked_date  
FROM inventory  
WHERE stock_quantity < 50  
ORDER BY last_stocked_date DESC
```

The following SQL statement incorrectly applies the minimum stock quantity filter at the category level, not the category and last stocked date level.

```
SELECT category, last_stocked_date  
FROM inventory  
GROUP BY category HAVING MIN(stock_quantity) < 50  
ORDER BY last_stocked_date DESC
```

The following SQL statement only returns a single record, the very latest last_stocked_date.

```
SELECT category, last_stocked_date  
FROM inventory  
WHERE stock_quantity < 50  
ORDER BY last_stocked_date DESC LIMIT 1
```

[WHERE \(Transact-SQL\) - SQL Server | Microsoft Learn](#)

[GROUP BY \(Transact-SQL\) - SQL Server | Microsoft Learn](#)

Next >

[Check Your Answer](#)

Practice Assessment for Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric

Question 50 of 50

You have a Fabric lakehouse named Lakehouse1.

You write the following T-SQL statement targeting the SQL analytics endpoint of Lakehouse1.

```
SELECT CalendarYear  
, SalesAmount  
, **<target1>** (SalesAmount, **<target2>**, 0) OVER(ORDER BY CalendarYear) AS  
PreviousSalesAmount  
FROM dbo.Sales  
WHERE CalendarYear BETWEEN 2020 AND 2023
```

You need to compare the sales amount and ensure that the statement displays the value from the previous year in the PreviousSalesAmount column.

The following is an example of the expected result.

CalendarYear	SalesAmount	PreviousSalesAmount
2020	100	0
2021	200	100
2022	150	200
2023	500	150

Which function should you insert for <target1>, and which numeric value should you provide for <target2> in the statement?

LAG and 1

✓ This answer is correct.

LAG and -1

- LEAD and 1
- LEAD and -1

The `LAG()` function accesses the data from a previous row in the same result set by using a given physical offset. The offset argument represents the number of rows returned from the current row from which to obtain a value. Offset cannot be a negative value.

[LAG \(Transact-SQL\) - SQL Server | Microsoft Learn](#)

[Submit >](#)

[Check Your Answer](#)